

Cache Simulator Implementation Report

Pallav Kamad 2023CS51067 Shreeraj Jambhale 2023CS50048

Abstract

This report details the implementation of a cache simulator for a four-core system using the MESI coherence protocol. The simulator models private L1 caches with write-back, write-allocate policies and LRU replacement. Key features include bus arbitration, snooping-based coherence, and detailed performance statistics. This document explains the high-level design, core classes, data structures, and critical workflows, providing insight into how memory operations are processed, synchronized, and optimized.

Contents

1	Implementation Overview	2
1.1	Simulator Components	2
1.2	Key Design Choices	2
2	Main Classes and Responsibilities	2
2.1	CacheSimulator Class	2
2.2	L1Cache Class	2
3	Core Data Structures	3
3.1	CacheLine	3
3.2	cache_sets	3
3.3	BusRequest	3
3.4	trace_data	3
4	Key Function Workflows	3
4.1	CacheSimulator::runSimulation()	3
5	Performance Evaluation	4
5.1	Metrics Collected	4
5.2	Sample Results	4
5.3	Optimizations Observed	4
5.4	Experimental Section	4
5.5	Experimental Section 2	5
6	Limitations and Future Work	6
6.1	Current Limitations	6
7	Conclusion	6

1 Implementation Overview

1.1 Simulator Components

- **CacheSimulator:** Manages global simulation state, bus arbitration, and trace loading.
- **L1Cache:** Represents each core's private cache, handling hits/misses, coherence actions, and LRU replacement.

1.2 Key Design Choices

- **MESI Protocol:** Ensures coherence via Modified, Exclusive, Shared, and Invalid states.
- **Central Snooping Bus:** Serializes transactions using a FIFO queue with priority for the lowest core ID.
- **Cycle-Accurate Modeling:** Tracks execution and idle cycles for each core, simulating blocking behavior during misses.
- **Statistics Collection:** Logs hits, misses, evictions, writebacks, and bus traffic for analysis.

2 Main Classes and Responsibilities

2.1 CacheSimulator Class

Roles:

- Loads memory traces from files into per-core trace_data.
- Manages the global clock (cycle), bus queue (bus_queue), and arbitration logic.
- Coordinates core execution and bus transactions.

Critical Methods:

- `runSimulation()`: Drives the simulation loop, advancing cores and resolving bus requests.
- `loadTraces()`: Parses input trace files into memory references.

2.2 L1Cache Class

Roles:

- Processes load/store requests, triggering hits, misses, or coherence operations.
- Implements LRU replacement and MESI state transitions.
- Snoops bus transactions from other cores.

Critical Methods:

- `processMemoryRequest()`: Handles hits/misses, enqueues bus requests (BUS_RD, BUS_RDX, BUS_UPGR).
- `handleBusRequest()`: Downgrades/invalidates lines during snooping.
- `completeMemoryRequest()`: Installs new cache lines after bus transactions.

3 Core Data Structures

3.1 CacheLine

```
struct CacheLine {
    bool valid;
    uint32_t tag;
    MESIState state;
    int lru_counter;
    std::vector<uint8_t> data; // Block data (size = 2^b bytes)
};
```

Tracks validity, coherence state, and LRU metadata for each cache block.

3.2 cache_sets

- A 2D vector: `std::vector<std::vector<CacheLine>> cache_sets`.
- Organized as **S sets** \times **E ways** (configurable via s and E parameters).

3.3 BusRequest

```
struct BusRequest {
    int core_id;
    BusOperation operation; // BUS_RD, BUS_RDX, BUS_UPGR, FLUSH
    uint32_t address;
    int start_cycle;
    int duration;
};
```

Represents pending coherence transactions (e.g., read, invalidate, upgrade).

3.4 trace_data

- Stores memory references for each core: `std::vector<std::vector<MemRef>> trace_data`.
- MemRef includes operation type (read/write) and address.

4 Key Function Workflows

4.1 CacheSimulator::runSimulation()

Flowchart:

1. **While** traces are not exhausted:
 - **Check Bus:** If free, dequeue the next BusRequest (lowest core ID first).
 - **Advance Cores:** For each unblocked core, process the next memory reference.
 - **Hit:** Update LRU, increment stats.
 - **Miss:** Enqueue bus requests (e.g., BUS_RD + FLUSH if evicting Modified line).
 - **Complete Bus Transactions:** Unblock cores and update cache states after duration cycles.
 - **Increment Global Clock.**

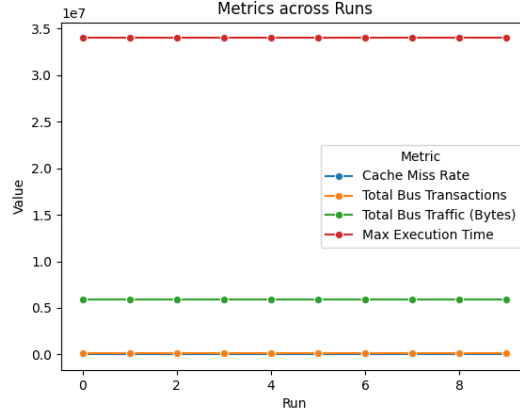


Figure 2: Output consistency across 10 simulation runs

5.5 Experimental Section 2

Running on different parameters to evaluate performance.

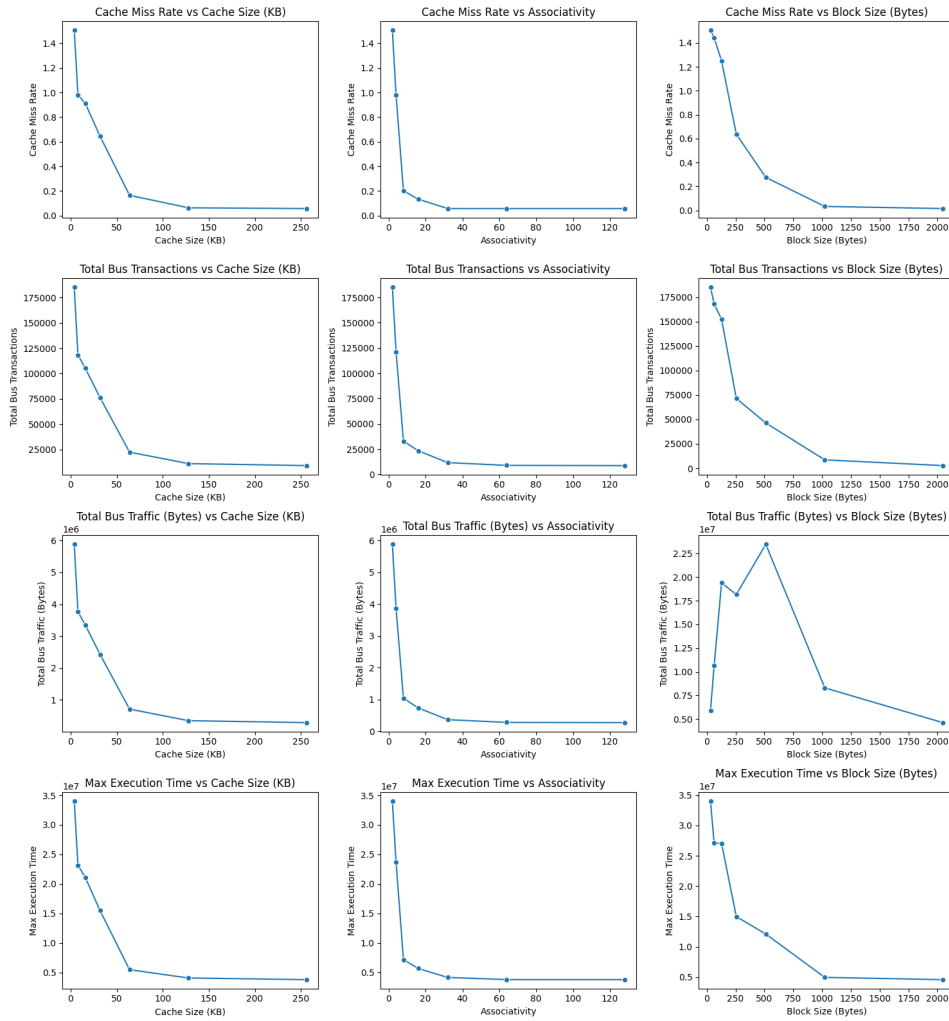


Figure 3: Performance on app1

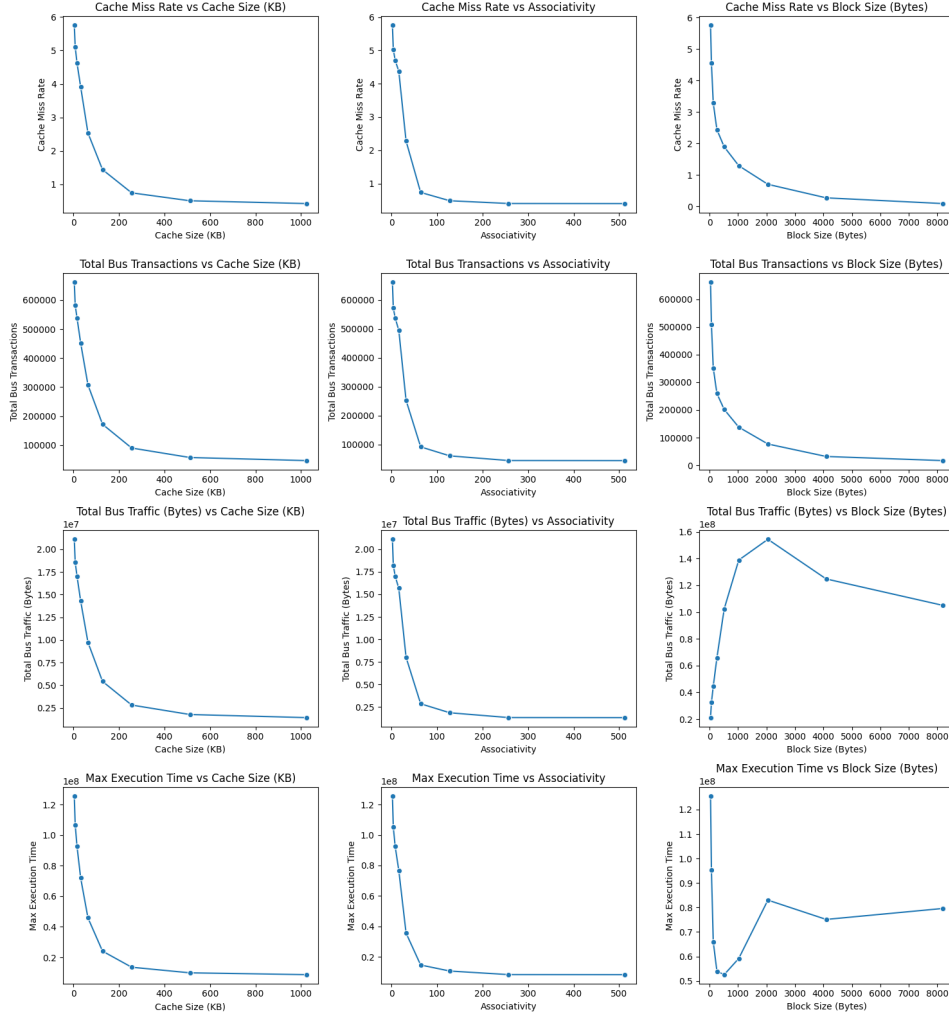


Figure 4: Performance on app2

6 Limitations and Future Work

6.1 Current Limitations

1. **Single-Level Caches:** No L2/L3 hierarchy.
2. **Blocking Design:** Cores stall during misses.
3. **Simplified Bus Model:** No pipelining or priority-based arbitration.

7 Conclusion

This simulator provides a modular, cycle-accurate model of a multicore cache system, emphasizing MESI coherence and LRU replacement. By tracking per-core and bus-level metrics, it enables analysis of spatial/temporal locality, invalidation overheads, and contention. The implementation's clarity and extensibility make it a foundation for exploring advanced features like nonblocking caches, prefetching, and heterogeneous architectures.

MESI Protocol State Transitions with Cycle Costs

Processor-Initiated Transitions

Transition	Trigger	Process	Cycles	Bus Usage
I → E	Read miss, no peer	Issue BUS_RD , fetch from memory, install Exclusive , record hit.	100+1=101	Bus busy for 100 cycles.
I → S	Read miss, peer supplies	Issue BUS_RD , cache-to-cache transfer, install Shared .	2N+1	Bus busy during transfer.
I → M	Write miss	Issue BUS_RDX , invalidate peers, fetch memory, install Modified .	100+1=101	Bus busy for 100 cycles.
S → M	Write hit on Shared	Issue BUS_UPGR , invalidate peers, upgrade to Modified .	1	Instant invalidation.
E → M	Write hit on Exclusive	Local state change to Modified .	1	None.
Read hit	Read in S/E/M	Return data from cache.	1	None.
Write hit	Write in M	Local update in cache.	1	None.

Table 1: Processor-Initiated MESI Transitions

Bus-Observed Transitions (Snooping)

Transition	Trigger	Process	Cycles	Core Stall
M → S	BUS_RD from peer	Downgrade to Shared , supply data cache-to-cache.	2N	No stall.
M → I	BUS_RDX from peer	Write back, invalidate local copy.	100	No stall.
E → S	BUS_RD from peer	Downgrade to Shared , supply data.	2N	No stall.
E → I	BUS_RDX from peer	Invalidate local Exclusive .	2N	No stall.
S → I	BUS_RDX / BUS_UPGR	Invalidate local Shared .	0	No stall.

Table 2: Bus-Observed MESI Transitions